

Advanced Operating Systems (263-3800-001) Caches and TLBs

Timothy Roscoe and Andrew Baumann

Based on slides by Kevin Elphinstone & Gernot Heiser (UNSW)



© Systems Group | Department of Computer Science | ETH Zürich

Thursday 14 October 2010

Caching (review)

- Cache: something that remembers previous results
 - ⇒ can *sometimes* give the answer faster
 - E.g. memory caches, TLBs, etc.
- Work by *locality*:
 - Temporal locality:
if I needed x recently, I'm likely to need it again soon.
 - Spatial locality:
if I needed x, I'm likely to need something close by.

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Overview


- Review of caches
 - Assume you know: direct mapping, associativity, ...
- Cache addressing schemes
- Synonyms and Homonyms
- Cache management in Operating Systems
- Translation Lookaside Buffers
 - Coverage
- ARM cache/MMU architecture

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Memory caching (review)

- Fast memory between registers & slow RAM
 - 1–5 vs. 10–100 cycles
- Holds recently used data and/or instructions
- Compensates for slow RAM if *hit rate* high (~90%)
- Hardware: (mostly) transparent to software
- Size: few kB – several MB
- Typically hierarchy (2 – 5 levels)
 - on- & off-chip



Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science


Cache organization

- Transfer units
 - registers ↔ L1 cache ≤ 1 word (1 – 16 bytes)
 - cache ↔ RAM (or cache) 16 – 32 bytes, or more
 - Cache line*: also unit of storage allocation in cache
- Cache line associated information:
 - Valid bit
 - Modified bit
 - Tag
- Improves memory access by:
 - Absorbs reads (increases b/w, reduces latency)
 - Make writes asynchronous (hides latency)
 - Clusters reads & writes (hides latency)

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Cache access

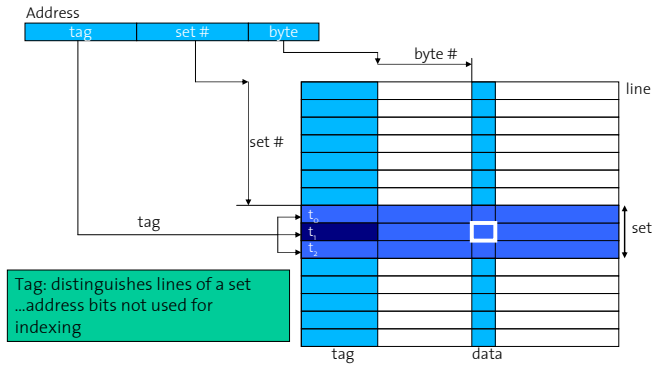


- Virtually indexed*:
 - Lookup by virtual address
 - Concurrent with address translation
- Physically indexed*:
 - Lookup by physical address

Thursday 14 October 2010


Advanced Operating Systems | Systems Group | Department of Computer Science

Cache indexing



Cache mapping

- Different memory locations map to same cache line



- Locations mapping to set i are said to be of colour i
- n -way assoc. caches hold n lines of the same colour
- Cache miss types:
 - Compulsory miss: data cannot be in cache
 - Capacity miss: all cache entries in use by other data
 - Conflict miss: set mapped to address is full
 - Can't happen in fully-associative cache
 - Coherence miss: forced by hardware coherence protocol

Cache addressing

- Address hashed \Rightarrow index of *line set*
- Associative lookup within set using *tag*
- n lines/set: *n-way set-associative cache*
 - $n=1$: direct mapped
 - $n=\infty$: fully associative
 - Usually $n=1-4$, occasionally $32-64$
- Hashing must be simple \Rightarrow
 - Use least significant bits of address

Cache write policy

- Write back: store only updates cache
 - Memory updated when dirty line replaced (flushed)
 - Clusters writes
 - Memory inconsistent with cache
 - Hard for multiprocessors
- Write through: store updates cache & memory
 - Memory always consistent with cache
 - Increased memory/bus traffic
- Store to line not currently in cache:
 - Write allocate: allocate new cache line & store
 - No allocate: store to memory & bypass cache
- Typical combinations:
 - Write-back & write-allocate
 - Write-through & no-allocate

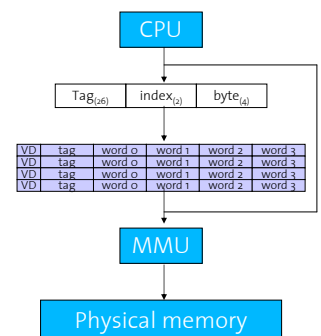
Cache addressing schemes

- So far, assumed cache *only* sees a virtual or physical address
- But indexing and tagging can use different addresses
 - Virtually-indexed, virtually tagged (VV)
 - Virtually-indexed, physically-tagged (VP)
 - Physically-indexed, virtually tagged (PV)
 - Physically-indexed, physically-tagged (PP)

Complex & rare

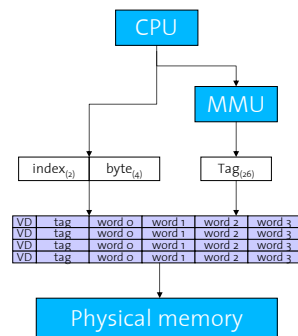
Virtually indexed, virtually tagged

- Also called:
 - Virtually addressed
- Also (misleadingly!)
 - Virtual cache
 - Virtual address cache
- Only uses virtual addresses
 - Operates concurrently with MMU
- Often used on-core




Virtually indexed, physically tagged

- Virtual addr \Rightarrow line
- Physical addr \Rightarrow tag
- Address translation required to retrieve data
- Index concurrently with MMU
- Use MMU output for check
- Typically used on-core



Physically indexed, physically tagged

- Only uses physical addresses
- Translation must complete before cache access can start
- Typically used off-core
- Note: page offset invariant under virtual address translation
 - Index bits \subseteq offset
 - Cache accessed without translation
 - Can be used on-core



Cache issues

- Caches are hardware, transparent to software
- So, why worry about them in the OS?
- Well...
 - Performance
 - Synonyms
 - Homonyms
- And pretty much essential for **multiprocessors**
 - Can't really scale without them
 - A later lecture will cover MP/OS cache issues

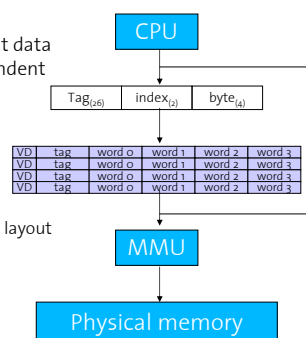
Cache performance really matters.

	Cycles	Normalised
L1 cache	2	1
L2 cache	15	7.5
L3 cache	75	37.5
Other L1/L2	130	65
Memory	~300	~150
1-hop remote L3	190	95
2-hop remote L3	260	130

AMD "Barcelona" @ 2GHz

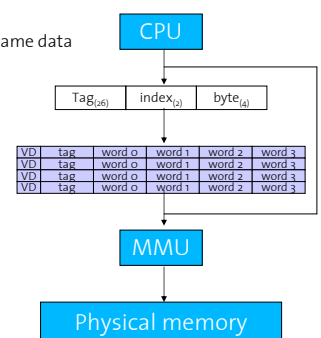
Virtually-indexed cache issues

- Homonyms**: same names for different data
- VA used for indexing is context dependent
 - Same VA refers to different PAs
 - Tag does not uniquely identify data
 - Wrong data is accessed!
 - OS must prevent this!
- Homonym prevention:
 - Flush cache on context switch
 - Force non-overlapping address-space layout
 - Tag** VA with **address-space ID** (ASID)
 - Use physical tags




Virtually-indexed cache issues

- Synonyms** (aliases): multiple names for same data
 - Several VA map to the same PA
 - Frames shared between processes
 - Multiple mappings of frame within AS
- May access stale data:
 - Same data cached in several lines
 - On write, one synonym updated
 - Read on other synonym returns old value
 - Physical tags don't help!
 - ASIDs don't help!
- Are synonyms a problem?
 - Depends on page and cache size
 - No problem for r/o data
 - Or i-caches



Example: MIPS R4x00 synonyms

- ASID-tagged, on-chip L1 VP cache
 - 16kB cache, 32B lines, 2-way set associative
 - 4kB (base) page size
 - Set size = $16\text{kB}/2 = 8\text{kB} > \text{page size}$
 - Overlap of tag & index bits, but from different addresses!




- Remember, location of data in cache determined by index
 - Tag only confirms whether it's a hit
 - Synonym problem iff $VA_{12} \neq VA'_{12}$
 - Similar issues on other processors, e.g. ARM11 (set size 16kB, page size 4kB)

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Address mismatch problems: Aliasing




- Page aliased in different address spaces
 - AS1: $VA_{12} = 1$, AS2: $VA_{12} = 0$
- One alias gets modified
 - In a write-back cache, other alias sees stale data
 - Lost-update problem

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Address mismatch problems: Remapping




- Unmap page with dirty cache line
- Re-use (remap) frame to a different page (in same or different AS)
- Write to a new page
 - Without mismatch, new write overwrites old (hits same cache line)
 - With mismatch, order can be reversed: "cache bomb"

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

DMA consistency problem



- DMA (normally) uses physical addresses and bypasses cache
 - CPU access inconsistent with device access
 - Need to flush cache before device write
 - Need to invalidate cache before device read

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Avoiding synonym problems

- Hardware synonym detection
- Flush cache on context switch
 - Doesn't help for aliasing *within* address space
- Detect synonyms and ensure
 - All read only, OR
 - Only one synonym mapped at a time
- Restrict VM mapping so synonyms map to same cache set
 - E.g. on R4x00, ensure that $VA_{12} = PA_{12}$

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Summary: VV caches

- Fastest: don't rely on TLB for retrieving data
 - Still need TLB lookup for protection
 - Or other mechanism to provide protection
- Suffer from synonyms and homonyms
 - Requires flush on context switch
 - Makes context switches expensive
 - May even be required on kernel \Rightarrow user switch
 - ... or guarantee of no synonyms or homonyms
- Require TLB lookup for writeback!
- Used on i860, ARM7/ARM9/StrongARM/XScale
- Used for i-cache on many architectures
 - Alpha, Pentium 4, etc.

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Summary: VV caches with keys

- Add *address space identifier* (ASID) part of tag
- On access compare with CPU's ASID register
- Removes homonyms, creates synonyms
 - Potentially better context switching performance
 - ASID recycling still requires a cache flush
- Doesn't solve synonym problem
 - (but that's less serious)
- Doesn't solve write-back problem!

VP caches


- Medium speed
 - Lookup in parallel with address space translation
 - Tag comparison after address translation
- No homonym problem
- Potential synonym problem
- Bigger tags (can't leave off set-number bits)
 - Increases area, latency, power consumption
- Used on most modern L1 data caches

PP caches

- Slowest
 - Requires result of address translation before lookup starts
- No synonym problem
- No homonym problem
- Easy to manage
- If small or highly associative (all index bits come from page offset), indexing can be in parallel with address translation
 - Potentially useful for L1 cache, e.g. Itanium
- Cache can use *bus snooping* to receive/supply DMA data
- Usable as off-chip cache with any architecture

Write buffers

- Store operations take long to complete
 - E.g. if cache line must be read or allocated
- Can avoid stalling CPU by buffering writes
- **Write buffer** is FIFO queue of incomplete stores
 - also called **store buffer** or **write-behind buffer**
- Can also read intermediate values out of buffer
 - To service load of a value that is still in the write buffer
 - Avoids unnecessary stalls of load operations
- Implies that memory contents are temporarily stale
 - On a multiprocessor, CPUs see different order of writes
 - "weak store order", to be revisited in SMP context!




Summary

- The OS has to *manage caches* if it is to provide:
 - Correctness
 - Performance
- Interactions between caches and memory translation are complex and subtle
- OSes typically try to hide these from the user

Translations Lookaside Buffer (TLB)

- TLB is a (VV) cache for page table entries
- TLB can be
 - Hardware loaded, transparent to OS, or
 - Software loaded, maintained by OS
- TLB can be:
 - Split, instruction and data TLBs, or
 - Unified
- Modern, high performance architectures use a hierarchy of TLBs:
 - Top-level TLB hardware-loaded
 - from lower levels
 - Transparent to OS



TLB issues: Associativity

- First TLB (VAX 11/780) was 2-way associative
- Most modern architectures have fully associative TLBs
- Exceptions:
 - i486, Pentium, P6, ... (4-way)
 - IBM RS/6000 (2-way)
- Reasons:
 - Modern architectures support multiple page sizes (superpages)
 - Better utilises TLB entries
 - TLB lookup done without knowing page base address
 - Set-associativity loses speed advantage
 - Hence superpage TLBs are fully-associative

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

TLB Coverage:

- Memory sizes are increasing
- # TLB entries more-or-less constant
- Pages sizes are growing *very* slowly
 - Total RAM mapped by TLB is not changing much
 - Fraction of RAM mapped by TLB is shrinking lots
- Modern architectures have very low *TLB coverage*
- Also, some architectures have software-loaded TLBs
 - General increase in TLB miss penalty (handling cost)
- The TLB is becoming a performance bottleneck

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

Concrete example: ARM

- Typical features of older ARM MMU cores:
 - Virtually-indexed L1 split caches
 - No L2 cache
 - No address-space tags (ASIDs) in TLB or caches
- Warning:
 - Under some circumstances, L4 does not flush caches on a context switch.
 - Instead, uses *domain bits*

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

TLB Size (i-TLB + d-TLB)

Architecture	TLB Size	Page Size	TLB Coverage
VAX	64-256	512B	32-128kB
ia32 / x86 (typical)	32-32+64	4kB+4MB	128-128+256kB
MIPS	96-128	4kB-16MB	384kB - ...
SPARC	64	8kB-4MB	512kB - ...
Alpha	32-128+128	8kB-4MB	256kB - ...
RS/6000	32+128	4kB	128+512kB
Power4/G5	128	4kB+16MB	512kB - ...
PA-8000	96+96	4kB-64MB	
Itanium	64+96	4kB-4GB	

Not grown much in 20 years!

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science


Address space usage vs. TLB coverage

- Each TLB entry maps 1 virtual page
- On TLB miss, reloaded from page table (PT), which is in memory
 - Some TLB entries needed to map page tables
 - Eg. 32-bit page table entries, 4kB pages.
 - One PT page maps 4MB
- Traditional UNIX process has 2 regions of allocated virtual address space
 - Low end: text, data, heap
 - High end: stack
 - 2 – 3 PT pages are sufficient to map most address spaces
- Superpages can be used to extend TLB coverage
 - But difficult to manage in the OS

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

ARM (xscale) cache architecture




- Virtually-indexed caches:
 - flush caches on context switch
 - direct cost: 1k – 18k cycles
 - indirect cost: up to 54k cycles

Thursday 14 October 2010

Advanced Operating Systems | Systems Group | Department of Computer Science

ARM TLB format

- TLB has no PID tag
 - Must *flush* TLB on context switches
 - Direct cost: 1 cycle
 - Indirect cost: 3k cycles
 - TLB flush \Rightarrow cache flush
- Windows CE avoids flushing by
 - No protection!
 - Max 32 processes
- Better: use *domains*
 - Impose additional access restrictions
 - Simulate address space tags
 - Flush TLB lazily on collisions



ARM domains

- Every PTE is in a *domain*.
 - There are 16 in total.
- Each domain has a 2 bit field in the DACR specifying access
 - Access rights to many pages can be changed at once
 - Access faults \Rightarrow trap to kernel
- Exercise: (!)
 - Work out how to use this to share page tables between processes, and avoid most cache/TLB flushes on context switch

Summary

- OS Management of the TLB is critical for
 - Correctness
 - Performance
- Hardware is diverse
 - Many processors/MMUs have unusual features
 - Effective use of these requires thought and ingenuity
 - And isn't portable