

# Security

Advanced Operating Systems  
(263-3800-00)

Timothy Roscoe

Thursday 2 December 2010

# Overview



- Security is a huge field in its own right
- This lecture covers a selection of topics from the OS designer's perspective
  - There is a focus on security mechanisms
- ***Information Security*** course as background

# Outline

- Security introduction (from an OS perspective)
  - What does it mean for a system to be secure?
  - Trusted computing base
- Security mechanisms
  - Access control matrix
  - Access control lists (review)
  - Capabilities (in depth)
- Non-discretionary access control
  - Mandatory access control
- Decentralised information flow control (labels)

# Outline

- Security introduction (from an OS perspective)
  - What does it mean for a system to be secure?
  - Trusted computing base
- Security mechanisms
  - Access control matrix
  - Access control lists (review)
  - Capabilities (in depth)
- Non-discretionary access control
  - Mandatory access control
- Decentralised information flow control (labels)

# What does it mean for a system to be secure?



- We all have different ideas about what security means
- Saying that “a system is secure” is meaningless without specifying the **security policy**

## Definition:

A set of rules and practices that specify or regulate how a system or organization provides security services to protect sensitive and critical system resources. [RFC 2828]

- Policy specifies allowed states of the system
- **Security mechanisms** are used to enforce the policy

# Trusted computing base (TCB)



## Definition

The totality of protection mechanisms within a computer system, including hardware, firmware, and software, the combination of which is responsible for enforcing a security policy. [RFC 2828]

- That part of the system that must be relied upon to enforce a security policy
  - . . . or that can circumvent the security policy
- **Trusted** by definition,
  - but that doesn't make it **trustworthy**

# How can we make the TCB more trustworthy?



- Testing
- Source code inspection
  - . . . but how many bugs do you think are left in Linux?
- Assurance standards: orange book, common criteria
  - Various levels
  - Windows, Linux, and many others have been certified
- Formal verification
  - seL4 / L4.verified projects [Klein et al., 2009]
- **Make it smaller!**
  - Less code to trust
  - Less code to inspect/verify

# Outline

- Security introduction (from an OS perspective)
  - What does it mean for a system to be secure?
  - Trusted computing base
- Security mechanisms
  - Access control matrix
  - Access control lists (review)
  - Capabilities (in depth)
- Non-discretionary access control
  - Mandatory access control
- Decentralised information flow control (labels)



# Access control matrix

Representation/definition of permissible operations in a system:

	Objects				
Subjects	User1	User2	User3	File1	File2
User1		Send msg		RW	R
User2	Send msg				RW
User3	Set passwd	Set passwd	Set passwd		R

**Subjects:** users, processes, groups, etc.

**Objects:** other users/processes, files, memory objects, etc.

**Privileges/rights:** depends on object

(for file: read, write, execute, etc.)

# Access control matrix properties



- Dynamic data structure, frequent changes
- Very sparse with many repeated entries
- Impractical to store explicitly
- Most common discretionary mechanisms:
  - **Access control list**: stores a column (who can access this)
  - **Capabilities**: store a row (what can access)

# Issues for discretionary access control

- **Propagation**: Can a subject grant access to another?
- **Restriction**: Can a subject propagate a subset of its own rights?
- **Revocation**: Can access, once granted, be revoked?
- **Amplification**:  
Can an unprivileged subject perform restricted operations?
- **Determination of object accessibility**:  
Which subjects have access to a particular object?
  - Is an object accessible by any subject?  
(garbage collection)
- **Determination of subject's protection domain**:  
Which objects are accessible to a particular subject?

# Access control lists

- Implemented by most commodity systems
- **ACL** associated with **object**
  - Propagation: meta right (eg. owner may chmod)
  - Restriction: meta right
  - Revocation: meta right
  - Amplification: protected invocation right (eg. setuid)
  - Accessibility: explicit in ACL
  - Protection domain: hard (if not impossible) to determine
- Usually condensed via **groups / classes**
- Can have **negative rights**
- Sometimes implicit (eg. UNIX process hierarchy)

# UNIX ACLs

- Despite modern terminology, classic UNIX privileges are a (restricted) ACL representation:

drwx--x--x	2	andrewb	andrewb	4.0K	Mar	9	22:28	dir1
-rwxr-xr-x	1	andrewb	andrewb	1.3K	Mar	9	22:26	file1
-rw-r--r--	1	andrewb	andrewb	13K	Mar	9	22:26	file2
-rw-----	1	andrewb	andrewb	30K	Mar	9	22:26	file3
-rw-rw----	1	andrewb	group1	92K	Mar	9	22:27	file4

Diagram illustrating the components of a classic UNIX file listing (ls -la) as a restricted ACL representation:

- rights**: Points to the permission string (e.g., drwx--x--x).
- subject (owner)**: Points to the owner field (e.g., andrewb).
- subject (group)**: Points to the group field (e.g., andrewb or group1).
- object**: Points to the file/directory name (e.g., dir1, file1, file2, file3, file4).

- Permissions for *other* are an implicit group of subjects

# Capabilities: introduction



- **Capability list** associated with **subject**
- Each capability confers a certain right to its holder
  - Propagation: copy/transfer capabilities between subjects
  - Restriction: requires creation of new (derived) caps
  - Revocation: requires invalidation of caps from all subjects
  - (may be difficult)
  - Amplification: special invocation capability
  - Accessibility: requires inspection of all capability lists
  - (hard if not impossible to determine)
  - Protection domain: explicit in capability list
- Few successful commercial systems:
  - IBM System/38 (AKA AS/400, iSeries, System i, . . . ), KeyKOS

# Capabilities

- Main advantage of capabilities is fine-grained access control
  - ☑ Easy to provide access to specific subjects
  - ☑ Easy to delegate permissions to others
    - See *The Confused Deputy*, Norm Hardy
- A cap presents prima facie evidence of the right to access
  - Think of it as a key
- Consists of **object identifier** and a set of **access rights**
  - Implies object naming
  - Any representation must protect capabilities against forgery
- How are caps implemented and protected?
  - Tagged**: protected by hardware
  - Partitioned**: protected by software
  - Sparse**: protected by sparsity  
(probabilistically secure, like encryption)

# Tagged capabilities, aka Hardware capabilities



Extra **tag bit** with every memory word (or group thereof)

- Tag identifies capabilities
- Capabilities may be used and copied like “normal” pointers
- Hardware checks permissions when dereferencing capability
- Modifications turn the tag off (reverting caps to plain data)
- Only the kernel can turn a tag bit on
- ☑ Propagation easy
- Restriction requires kernel to create new (weaker) capability
- ☒ **Revocation** virtually impossible (requires memory scan)
- ☒ **Accessibility** virtually impossible to determine



# Tagged capabilities outside RAM



- Disk has no tags
- AS/400 simulates them by restricting physical I/O to the low-level OS
  - Extra bit stored for every word on disk
  - On page-out, page must be scanned and tags collected
  - On page-in, tags are reconstructed
  - Significant processing overhead on all disk I/O

# Tagged capabilities: summary

- Secure through hardware protection
- Convenient for applications (appear as normal pointers)
- Checked by hardware  $\Rightarrow$  fast validation
- Capability hardware is complex (hence slow)
- Separate mechanisms required for I/O and distribution

# Partitioned capabilities



System maintains capabilities for each process, eg. as a **capability list (clist)**

- User code uses only handles (indirect references) to caps
- System validates access when performing any privileged operation (eg. mapping a page)
  - Validation is explicit at syscall time
  - Propagation: system call to copy a cap between clients
  - Restriction: invoke kernel to create new cap
  - Revocation: invoke kernel to remove cap from clist
  - Accessibility: requires scanning all clists
  - Protection domain: explicitly represented in clist
- Used in Hydra, Mach, KeyKOS, EROS, seL4, Barrelfish, many others


# Partitioned capabilities

- Secure through kernel protection
  - Real caps live only in kernel space
- Validation at mapping/invocation time
  - Apps use “normal” pointers
- ☑ Fast validation possible
- (for memory objects, validation is cached byMMU)
- Propagation requires kernel intervention
- ☑ Reference counting and revocation possible with kernel support
- ☑ No special hardware requirements

# Sparse capabilities

- Basic idea similar to encryption
- Add bit string to capability
  - Makes valid capabilities a tiny subset of the capability space
  - Secure by infeasibility of exhaustive search of cap space


# Encrypted sparse capabilities



Signature: bit string is encrypted object info

- Cap consists of object ID, rights, and signature
- **Signature** = object ID and rights encrypted with a private key
- Validated by checking signature

# Password sparse capabilities



Global  
object  
table

OID	Password	Rights

**Password:** bit string is random data

- Cap consists of object ID, and password
- Rights determined by looking up password in a global object table

# Sparse capabilities

- Sparse caps are regular user-level objects
- Can be passed like other data
  - Similar to tagged caps, but without hardware support
  - Validated at invocation time (either explicitly or implicitly)

## Issues:

- Full mediation requires extra support
  - See Mungi
- High amplification of leaked data
  - Problem with covert channels



# Aside: Physical memory management with caps

[Elkaduwe et al., 2008, Klein et al., 2009]



- Problem: allocation and management of physical memory
- Most kernels use dynamic allocation (malloc) for metadata
  - What do you do when it runs out?
- seL4 model (also used in Barrelfish):
  - All physical memory manipulated as partitioned capabilities
  - Memory not used for bootstrap is initially *untyped*
  - User controls allocation by *retyping* capabilities to:
    - Frame** may be mapped into a page table
    - Thread** kernel TCB and metadata for a new thread
    - Endpoint** targets for IPC
    - VNode** hardware page tables
    - CNode** tables used to store further capabilities
  - No dynamic allocation necessary in kernel!

# Outline

- Security introduction (from an OS perspective)
  - What does it mean for a system to be secure?
  - Trusted computing base
- Security mechanisms
  - Access control matrix
  - Access control lists (review)
  - Capabilities (in depth)
- Non-discretionary access control
  - Mandatory access control
- Decentralised information flow control (labels)

# Non-discretionary access control

- Both ACLs and capabilities are *discretionary access control*
  - Discretionary, a user with access to data can pass it on

## Alternatives:

- Mandatory (system-controlled) policies
  - eg. certain users never access certain objects
  - No user can change these
  - Focus on restricting information flow
- Role-based policies
  - Subjects can take on specific pre-defined roles
  - Access rights depend on role

# Mandatory access control (MAC)

Example: Bell-LaPadula model

- Every object  $o$  has a **security classification**  $L(o)$
- Every subject  $s$  has a **security clearance**  $L(s)$
- Classifications & clearances form hierarchical **security levels**
  - eg. top secret > secret > confidential > unclassified
- Rule 1 (*no read up*)
  - $s$  can **read**  $o$  only if  $L(s) \geq L(o)$
  - standard confidentiality
- Rule 2 (*\* property*)
  - $s$  can **write**  $o$  only if  $L(s) \leq L(o)$
  - prevents leakage (accidental or intentional)
  - problems: logging, command chain
  - need way to declassify data

# Outline

- Security introduction (from an OS perspective)
  - What does it mean for a system to be secure?
  - Trusted computing base
- Security mechanisms
  - Access control matrix
  - Access control lists (review)
  - Capabilities (in depth)
- Non-discretionary access control
  - Mandatory access control
- Decentralised information flow control (labels)

# Problems with MAC

- ☑ Under MAC, you don't need to trust those who access data not to leak it
- ☒ However, only the system administrator may declassify data
  - Not very flexible; leads to data becoming more classified
- ☒ Single system-wide policy
  - Small pre-defined number of categories
- **Decentralized information flow control:**
  - Allows a large (essentially unlimited) number of categories
  - Extends MAC with the notion of an owner for each category
  - The owner may declassify data within that category
  - Data may belong to multiple categories

# Motivation



Why do I care about information flow control?

- **Asbestos**: you don't want to trust a dynamic web service implementation not to leak private data  
[Efsthopoulos et al., 2005]
- **HiStar**: you don't want to trust your virus scanner, login daemon, VPN client, ... [Zeldovich et al., 2006]
- Both systems implement essentially the same model
  - We'll use the HiStar terminology/notation

# Labels

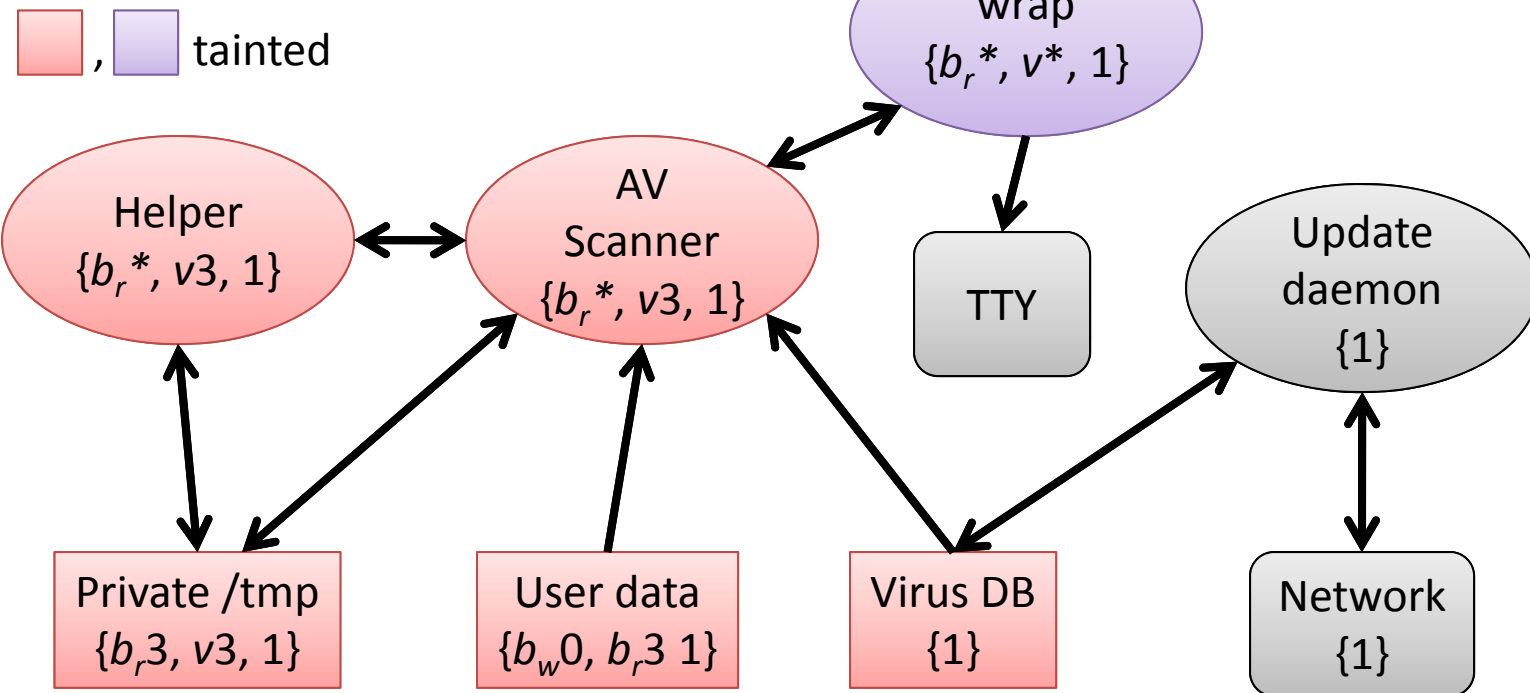
- Every object has a label
- Label specifies how tainted the object is *for every category*
  - Categories are 61-bit identifiers, may be freely created
- The label of a subject acquires **taint** based on its activities (eg. reading data)
- MAC checks apply to every operation for every category
  - An action is disallowed if it would convey information from more to less tainted objects in any category
- There are 5 **taint levels**:
  - \* has untainting privileges in this category
  - 0 cannot be written/modified by default
  - 1 default level – no restrictions
  - 2 cannot be untainted/exported by default
  - 3 cannot be read/observed by default



# Star level

- Subjects (threads and gates in HiStar) with the star (\*) level in a category are trusted for that category
  - Can disregard MAC rules (i.e. declassify)  
*but only within that category*
  - Called the **owner** of a category
- Essential difference from classic (military) MAC systems

# Untrusted virus scanner on HiStar



- wrap trusted to untaint scan results and display them
- $\{b_r^*; v3; 1\}$ : helper owns the  $b_r$  category, has level 3 in the v category, and level 1 otherwise
- $b_r$  = user's read category,  $b_w$  = user's write category

# Challenges for DIFC

- Efficient implementation of labels and label comparison
  - Time and memory
- Usability issues
  - How do you pick the labels?
  - How do you debug taints?  
Once an operation fails it's too late!

# Summary

- Access control mechanisms are provided by the OS to enforce a security policy
- Common discretionary mechanisms: ACLs and capabilities
  - Capabilities allow fine-grained access control
  - Caps enable simple and controlled delegation of privilege
- MAC provides information flow control, but is unsuited to a general-purpose system
- Decentralised information flow control (Asbestos, HiStar) tries to address this

# Summary

Other interesting topics we haven't covered:

- Role-based access control
- Principle of least authority/privilege
- Trusted computing / TPM
- Covert channels
- Formal verification

# References

- Efstathopoulos, P., Krohn, M., VanDeBogart, S., Frey, C., Ziegler, D., Kohler, E., Mazières, D., Kaashoek, F., and Morris, R. (2005). **Labels and event processes in the Asbestos operating system**. In 20th SOSP, pages 17–30.
- Elkaduwe, D., Derrin, P., and Elphinstone, K. (2008). **Kernel design for isolation and assurance of physical memory**. In IIES '08: Proc. 1st workshop on isolation and integration in embedded systems, pages 35–40.
- Klein, G., Elphinstone, K., Heiser, G., Andronick, J., Cock, D., Derrin, P., Elkaduwe, D., Engelhardt, K., Kolanski, R., Norrish, M., Sewell, T., Tuch, H., and Winwood, S. (2009). **seL4: Formal verification of an OS kernel**. In Proceedings of the 22nd ACM Symposium on Operating Systems Principles, pages 207–220.
- Zeldovich, N., Boyd-Wickizer, S., Kohler, E., and Mazières, D. (2006). **Making information flow explicit in HiStar**. In 7th OSDI.