



Domain-Specific Languages

Advanced Operating Systems
(263-3800-00)

Timothy Roscoe

Thursday 9th December 2010

Outline

- Introduction
- Interface definition languages
- Hardware interface languages
- Filet-of-Fish
- Hamlet
- References



The problem

- C is a pain to write OS code in.
- 2 classes of problem:
 1. Lack of automatic resource mgmt
 2. Hard to express high-level semantics



High-level languages to the rescue?

- Write your OS in Java/Eiffel/C#/etc.
 - Has been tried. Several times.
- Problems:
 - Lose all control over resource management
 - Explicit layout / memory access becomes hard
 - Still can't express high-level semantics
 - (OS code is highly specialized)
 - Sufficiently-expressive languages too slow and too abstract
 - (e.g. Haskell)



Extend C?

Promising approach:

- NesC: TinyOS's C dialect with support for modules, events [Gay 2003]
- Deputy: extensions to C using type inference for static checks [e.g. Anderson 2009]
- Ivy: evolving C as a language [Brewer 2005]

So far, little uptake (poor toolchain support?)



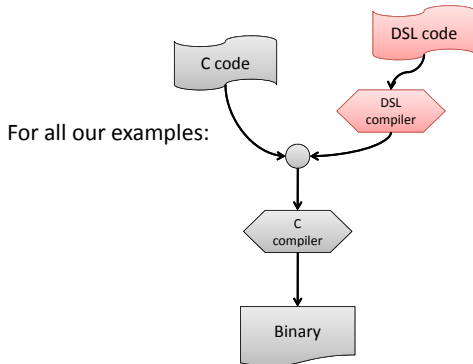
Domain specific languages

- Old idea
- Very broad applicability (not just OSes)
- Guy Steele: "design your system as if you were designing a language anyway".
- Build a "little language" tailored for the task at hand
- Generate C which is then compiled with the OS

In Barrellfish, we use DSLs extensively (4 so far, and counting)



Domain specific language workflow



Advantages



- Highly specialized: capture the exact semantics you want!
- Can check and enforce useful invariants
- Small, easy to learn
- Can be very fast (faster than a programmer could write)
- Dramatically reduces devel/debug time

Of course, there is a downside:

- Lot of effort to write the compiler
- Complicates toolchain management
- May make the code look somewhat alien...

Examples of DSLs in Operating Systems



- Communication interface definition
- Hardware register access
- Scheduling algorithms (see next week!)
- Protocol stack design (Click, Prolac)
- Capability type system specification
- Error code definitions
- ...

Outline



- Introduction
- Interface definition languages
- Hardware interface languages
- Filet-of-Fish
- Hamlet
- References

Interface definition languages



- Perhaps oldest DSLs for OS development
- Original RPC [Birrell and Nelson, 1984]

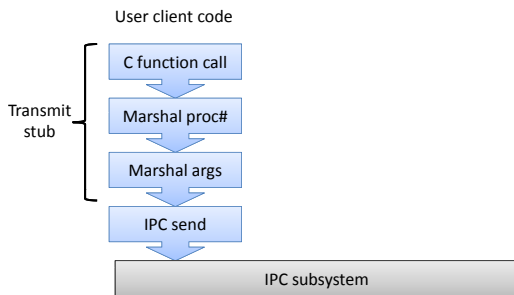
Interface definition semantics



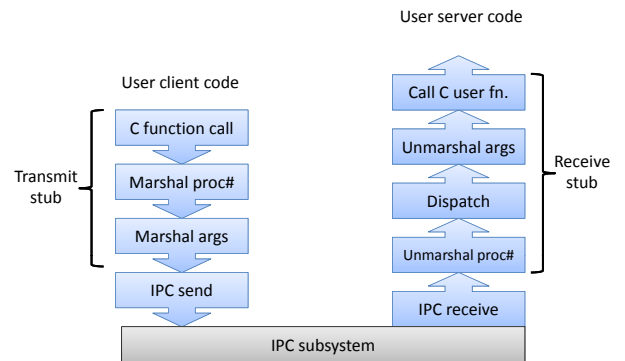
IDLs are NOT like (Java) RMI!

- An IDL typically defines its own type system.
- Concrete types: integers, structs, etc.
- Abstract types: interface references
- IDL compiler maps this to (perhaps many) programming languages

Stub functionality



Stub functionality



Memory management in IDLs



Network RPC IDLs need to worry about memory management:

- E.g. CORBA defines 3 parameter modes:
 - in**: Argument is passed from client to server (parameter)
 - out**: Argument is passed from server to client (result)
 - inout**: Argument is sent to server, modified, sent back

Memory management in IDLs



Network RPC IDLs need to worry about memory management:

- E.g. CORBA defines 3 parameter modes:
 - in**: Argument is passed from client to server (parameter)
 - out**: Argument is passed from server to client (result)
 - inout**: Argument is sent to server, modified, sent back
- This is not enough locally, in an OS with:
 - Shared memory transport
 - No garbage collection
 - Values (like arrays) bigger than a register or machine word

Memory management in IDLs



Basic questions:

- When should memory in the client be freed?
- How can memory in the server be allocated?
- How can memory in the client be allocated?
- When should memory in the server be freed?
- When is it safe to modify client data, if it's been sent to the server?

Memory management in IDLs

[Hamilton and Kougouris, 1994]



The IDL for the Spring OS modified CORBA IDL for an OS setting:

- copy**: Argument is copied to the server.
- consume**: Argument is sent from client to server, and destroyed at client.
- produce**: Argument is generated at the server and sent back (destroyed at server)
- borrow**: Like **inout**, but can't be modified by client in the meantime.

Performance



For network IDLs (CORBA, ANSA, DCE, SunRPC, etc.)
stub performance not critical

- Network latency dominates
- Calls are infrequent
- Calls must traverse network stack anyway

Performance



For network IDLs (CORBA, ANSA, DCE, SunRPC, etc.)
stub performance not critical

- Network latency dominates
- Calls are infrequent
- Calls must traverse network stack anyway

It's very different for local OS stubs:

- IPC system highly optimized \Rightarrow stub performance critical
- Calls are frequent (particularly in a microkernel)

Flick

[Eide et al., 1997]



Optimizing stub compiler: many techniques, e.g.:

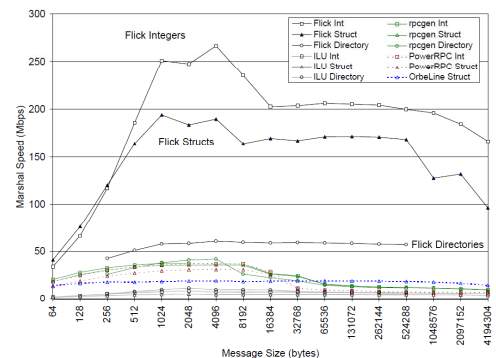
- Marshall all fixed-size data in one go
 - Avoid checking buffer size each time
- Inline most operations
- Use receive buffer space for arguments (e.g. in)
- Aggregate copies into one big `memcpy`
- Transport-specific marshalling
 - e.g. L4 IPC in registers

Flick performance

[Eide et al., 1997]



- Shows effect of compiler optimizations on marshalling code:



General pattern



Consider a domain specific language where:

- You're writing the same boilerplate code again and again (with minor variations)
 - It's easy to make mistakes
 - Interoperability (common specifications) are useful
 - It's clear what the compiler should do
 - Compiler optimizations would be useful
- ... or at least some of the above.

Outline



- Introduction
- Interface definition languages
- Hardware interface languages
- Filet-of-Fish
- Hamlet
- References

Hardware register access



Accessing hardware registers is generally fiddly code

- Lots of bit manipulation (registers have many fields)
- Poor C support
 - word size, sign extension, volatile semantics
 - bitfield structs are implementation specific!
- Consequences of errors are bad
 - Very hard to find bugs
 - Frequently hangs entire machine
- C code to manipulate registers is tedious to write

Devil Example: Logitech Busmouse

[Mérillon et al., 2000]



- Hand-written macros:

```
#define MSE_DATA_PORT      0x23c
#define MSE_CONTROL_PORT  0x23e
...
#define MSE_READ_Y_LOW     0xc0
#define MSE_READ_Y_HIGH    0xe0
```

Devil Example: Logitech Busmouse

[Mérillon et al., 2000]



- Hand-written macros:

```
#define MSE_DATA_PORT      0x23c
#define MSE_CONTROL_PORT  0x23e
...
#define MSE_READ_Y_LOW     0xc0
#define MSE_READ_Y_HIGH    0xe0
```

- Programmer usage idioms:

```
dy = (inb(MSE_DATA_PORT) & 0xf);
outb(MSE_READ_Y_HIGH, MSE_CONTROL_PORT);
buttons = inb(MSE_DATA_PORT);
dy |= (buttons & 0xf) << 4;
buttons = ((buttons >> 5) & 0x07);
```

Devil Example: Logitech Busmouse

[Mérillon et al., 2000]



- Device specified in the Devil DSL:

```
device logitech_busmouse (base : bit[8] port @ {0..3})
{
    // Signature register (SR)
    register sig_reg = base @ 1 : bit[8];
    variable signature = sig_reg, volatile, write trigger : int(8);

    // Configuration register (CR)
    register cr = write base @ 3, mask '1001000.' : bit[8];
    variable config = cr[0] : { CONFIGURATION => '1', DEFAULT_MODE => '0' };

    // Interrupt register
    register interrupt_reg = write base @ 2, mask '000.0000' : bit[8];
    variable interrupt = interrupt_reg[4] : { ENABLE => '0', DISABLE => '1' };

    // Index register
    register index_reg = write base @ 2, mask '1..00000' : bit[8];
    private variable index = index_reg[6..5] : int(2);

    register x_low = read base @ 0, pre {index = 0}, mask '.....' : bit[8];
    register x_high = read base @ 0, pre {index = 1}, mask '.....' : bit[8];
    register y_low = read base @ 0, pre {index = 2}, mask '.....' : bit[8];
    register y_high = read base @ 0, pre {index = 3}, mask '.....' : bit[8];

    structure mouse_state = {
        variable dx = x_high[3..0] # x_low[3..0], volatile : signed int(8);
        variable dy = y_high[3..0] # y_low[3..0], volatile : signed int(8);
        variable buttons = y_high[7..5], volatile : int(3);
    };
}
```

Devil Example: Logitech Busmouse

[Mérillon et al., 2000]



- What's Devil generating?

```
#define bm_get_mouse_state() ( \
    outb(1, bm_cache._dil_base_+2); bm_cache.cache_mouse_state.cache_get_x_high = inb(bm_cache._dil_base_); \
    outb(0, bm_cache._dil_base_+2); bm_cache.cache_mouse_state.cache_get_x_low = inb(bm_cache._dil_base_); \
    outb(3, bm_cache._dil_base_+2); bm_cache.cache_mouse_state.cache_get_y_high = inb(bm_cache._dil_base_); \
    outb(2, bm_cache._dil_base_+2); bm_cache.cache_mouse_state.cache_get_y_low = inb(bm_cache._dil_base_); \
)

#define bm_get_dy() ( \
    (bm_cache.cache_mouse_state.cache_get_y_high & 0xfu) << 4 | bm_cache.cache_mouse_state.cache_get_y_low & 0xfu)

#define bm_get_buttons() ((bm_cache.cache_mouse_state.cache_get_y_high & 0xe0u) >> 5)
```

Devil Example: Logitech Busmouse

[Mérillon et al., 2000]



- What the programmer gets to write:

```
bm_get_mouse_state();
dy = bm_get_dy();
buttons = bm_get_buttons();
```

Other Devil features



- Pre- and post-conditions
 - E.g. index registers used to access other register banks
 - Semaphores which must be held before writing a register
- “Variables”
 - values which combinations (usually concatenations) of register values

Mackerel features



- Goal: specifications should be as close to datasheet descriptions as possible.
- Basic constructs specify:
 - Individual **registers**
 - Register **types**
 - Register **arrays**
 - In-memory **data types**
 - Collections of **constant values**
- Make extensive of C compiler’s type system and inlining
- Comments are incorporated in C **printf**-like code

Mackerel



- Example: Intel e1000 Ethernet controller
- Fragment showing a register definition:

```
register status rw addr(base, 0x0008) "Device status" {  
    fd      1 "Link full duplex configuration";  
    lu      1 "Link up";  
    lan_id  2 "LAN ID";  
    txoff   1 "Transmission paused";  
    tbimode 1 "TBI mode";  
    speed   2 type(linkspeed) "Link speed setting";  
    asdv    2 type(linkspeed) "Auto speed detection val";  
    phyra   1 "PHY reset asserted";  
    -       8 mbz;  
    gio_mes 1 "GIO master enable status";  
    -       12;  
};
```

Mackerel

Barrelfish’s answer to Devil



- Things have changes somewhat in the meantime:
- Lots of address space ⇒ Index registers are less frequent
 - ⇒ pre-conditions less important
 - Register **address spaces** more useful (PCI, memory, IO)
- Registers are wider (32 or 64 bits)
 - ⇒ meaningful values rarely split across hardware fields
- Most complex devices communicate using **descriptor rings**
 - ⇒ In-memory data structures are just as important as registers

Mackerel features



Mackerel generates:

- C constant definitions for all constant values
- C Type definitions for all register and data types
- Functions to read/write all registers
- Functions to read/write all register and data type fields
- Functions to **snprintf**:
 - Register values
 - Data type values
 - Entire device state!

Mackerel: some figures



Lines of code (using David Wheeler’s SLOCCount):

2359 lines of Haskell for the Mackerel compiler
1028 lines of Mackerel for the e1000 specification
23762 lines of C generated from **e1000.dev**

If DSLs are so good...



How come we don't see more of them in OS research?

- Quite hard to design a good one
 - Except Mackerel, all the DSLs in Barrelfish were designed *after* we had an initial C implementation and understood the functionality.
- Perception: the effort to implement DSL usually outweighs the cost of designing, building, and implementing it
 - With yesterday's tools, there is some truth in this
 - But . . .

Building a DSL: what does it take?



DSLs are basically simple compilers:

Building a DSL: what does it take?



DSLs are basically simple compilers:

1. Parser
 - Used to be tedious to write
 - Glorious easy these days
 - E.g. combinator-based Monadic parsing in Haskell

Building a DSL: what does it take?



DSLs are basically simple compilers:

1. Parser
 - Used to be tedious to write
 - Glorious easy these days
 - E.g. combinator-based Monadic parsing in Haskell
2. Back-end C code generator
 - Rather more difficult . . .

Writing a backend for a DSL



The backend takes an AST and generates C code

- Basically: concatenate a set of strings into a C file
- Better: encode subset of C syntax into functional combinators easier.

But still:

- Writing code through a level of indirection
- Only captures syntax of C, not intended semantics.
- Can't automate tests
- Error-prone
- Annoying to debug
- Ultimately, no assurance it works.

Filet-o-Fish

[Dagand et al., 2009]



Filet-o-Fish is . . .

- Tool for writing C code generators
- Embedding of a subset of C in Haskell
- Notation for expressing DSL semantics
- Library for creating provably-correct C code from semantic specifications

Used in Barrelfish for (to date) 2 DSLs:

- Fugu** defines error codes and an error stack
- Hamlet** defines capability type system

Hamlet: specifying the capability type system

Yes, Hamlet really is a type of fish




How Filet-o-Fish compiles Hamlet

[Dagand et al., 2009]



Recall that Barrelfish uses typed, partitioned capabilities

- For each capability, we must specify:
 - Physical layout in memory
 - What it can be retyped to and from
 - Valid invocations on the capability
 - What happens when it is passed between domains
 - etc.
- We capture all this information in a Hamlet specification.



How Filet-o-Fish compiles Hamlet


[Dagand et al., 2009]



What does FoF look like?



- Defining semantics instead of syntax:



- For the previous example, Haskell resembles:

```
validateRetypeCode destType (srcType, validTypes) =
  do srcTypeV <- srcType
  validTypesP <- sequenceSem validTypes
  return (srcTypeV,
          (do returnc $ condition validTypesP))
where condition validTypes = fold orType false validTypes
      orType acc srcType = acc .||. (destType .==. srcType)
```

Using QuickCheck to test DSLs


[Dagand et al., 2009]



Affecting the OS design



- Check randomly-generated ASTs against semantic assertions:



Hamlet makes it easy to add new capability types to Barrelfish

- Led us to encode more functionality into the type system
 - E.g. different cap types for page table levels (On all architectures)
 - Type system enforces page table correctness
- Can encode multiple physical address spaces, etc.
- We expect to push further functionality into capability system...

Summary



Used appropriately:

- Reduce code complexity
 - Though rarely, if never, actually evaluated
 - DSLs perhaps seen more as a means to an end...
- Reduce bugs
 - Capture (and check) high-level semantics of the domain
- Facilitate automated testing and/or correctness proofs

References



- Birrell, A. D. and Nelson, B. J. (1984). **Implementing remote procedure calls**. ACM Trans. Comput. Syst., 2(1):39–59.
- Dagand, P.-E., Baumann, A., and Roscoe, T. (2009). **Filet-o-Fish: Practical and Dependable Domain-Specific Languages for OS Development**. In Proc. 5th Workshop on Programming Languages and Operating Systems (PLOS 2009).
- Eide, E., Frei, K., Ford, B., Lepreau, J., and Lindstrom, G. (1997). **Flick: A flexible, optimizing IDL compiler**. In PLDI, pages 44–56.
- Hamilton, G. and Kougiouris, P. (1994). **The Spring nucleus: A microkernel for objects**. Technical report, Sun Microsystems Laboratories.
- Mérillon, F., Réveillère, L., Consel, C., Marlet, R., and Muller, G. (2000). **Devil: An IDL for hardware programming**. In Proceedings of the 4th USENIX Symposium on Operating Systems Design and Implementation.

References



- David Gay, Philip Levis, Robert von Behren, Matt Welsh, Eric Brewer, and David Culler. 2003. **The nesC language: A holistic approach to networked embedded systems**. In Proceedings of the ACM SIGPLAN 2003 conference on Programming language design and implementation (PLDI '03). ACM, New York, NY, USA, 1-11
- Zachary R. Anderson, David Gay, and Mayur Naik. 2009. **Lightweight annotations for controlling sharing in concurrent data structures**. In Proceedings of the 2009 ACM SIGPLAN conference on Programming language design and implementation (PLDI '09)
- Eric Brewer, Jeremy Condit, Bill McCloskey, and Feng Zhou. 2005. **Thirty years is long enough: getting beyond C**. In Proceedings of the 10th conference on Hot Topics in Operating Systems - Volume 10 (HOTOS'05), Vol. 10. USENIX Association, Berkeley, CA, USA, 14-14.